

Calibration Impact on Comfort Scoring

Before / After Comparison on the Full 31-Recording Corpus

Heejung Roh Yiran Tao Sparsh Bansal Jung Yeop (Steve) Kim

This note reports the per-scenario impact of the calibration pipeline on the deployed comfort score. *Before* = hand-set default configuration (`config/default.yaml`); *After* = the calibrated configuration promoted to `config/deploy.yaml` after race 2 plus the no-pose ablation. Numbers are aggregated from a leave-one-recording-out (LORO) replay across all 31 RealSense `.bag` recordings; mean comfort and slope are computed over the intent + execution window per recording. Slope sign agreement is the fraction of recordings whose comfort slope matches the desired sign (`sc02`: positive; all others: negative).

Table 1: Per-scenario comfort metrics across all 31 LORO replays, before and after calibration. Numbers in **bold** highlight where calibration materially improved an entry against the desired behavior.

Scenario	n	Label	Mean comfort		Slope (pts/s)		Direction agree	
			Before	After	Before	After	Before	After
sc01_walkby	7	abort	69.2	50.2	-0.27	-0.39	5/7	5/7
sc02_comfortable	8	continue	68.3	92.7	-2.51	+0.73	0/8	7/8
sc03_gradual_disc.	5	abort	78.8	79.5	+0.16	-2.29	1/5	5/5
sc04_sudden_wdraw	5	abort	74.8	78.8	-0.96	-2.46	5/5	5/5
sc05_distracted	6	abort	73.2	73.0	-1.25	-3.28	6/6	6/6
Direction-correct total							17/31 (55%)	28/31 (90%)
Held-out test best Youden J ($n=8$)			0.00	1.00				
sc02 – sc01 mean gap			-0.9	+42.5				

Reading the table. Three things change under calibration. First, the *ordering* between abort and continue scenarios flips from inverted to a clean +42.5-point separation, driven entirely by `sc02` climbing from 68.3 to 92.7 and `sc01` falling from 69.2 to 50.2. Second, the *direction* of the comfort trajectory becomes correct on 28 of 31 recordings (90%) versus 17 of 31 (55%) under default — the largest gain is on `sc02`, where 0 of 8 successful handovers trended the right way under default, and 7 of 8 do under deploy. Third, `sc03` (gradual discomfort) flips from trending slightly *upward* (slope +0.16, wrong direction) to a strongly negative slope (-2.29), with direction agreement going from 1/5 to 5/5.

What didn't change. `sc04` (sudden withdrawal) and `sc05` (distracted) were already at 100% direction agreement under default; calibration steepens their downward slopes (more decisive aborts) but leaves the categorical classification unchanged. `sc01` (walk-by) direction agreement stays at 5/7 — the two outlier recordings remain ambiguous — but the absolute mean comfort drops 19

points, well below any plausible abort threshold, so the categorical decision is correct even on the disagreeing recordings.

The headline number. Best Youden’s J on the held-out test set ($n=8$, including 3 aborts) goes from 0.00 at any threshold $\tau \in [30, 80]$ under default to 1.00 at $\tau \in [55, 68]$ under deploy — the default configuration cannot separate abort from continue *at any threshold*, while deploy yields a wide margin window for the controller to operate in.

Optimization Search Space

Calibration tunes a 20-dimensional parameter vector via differential evolution against the cached parquets extracted in Stage A. Pretrained backbone weights (RetinaFace, EfficientNet-B0, L2CS-Net, MediaPipe Pose) are frozen; only the fusion / decision-layer parameters listed in Table 2 are searched. The abort threshold τ^* is selected post-hoc from the held-out ROC curve and is not part of the DE search. Source of truth: `PARAM_NAMES` and `BOUNDS` in `scripts/optimize_params.py`.

Table 2: Variables tweaked during few-shot optimization. Twenty parameters spanning fusion weights, sub-score coefficients, gaze/posture event thresholds and penalties, temporal smoothing, and missing-detection EMA decay. Bounds were chosen to keep DE candidates physically interpretable without artificially clipping the search.

Group	Parameter	Bounds	Role
Phase-aware fusion weights	<code>we_intent</code>	[0.1, 0.9]	emotion weight, intent phase
	<code>wp_intent</code>	[0.1, 0.9]	posture weight, intent phase
	<code>wg_intent</code>	[0.0, 1.0]	gaze weight, intent phase
	<code>we_exec</code>	[0.1, 0.9]	emotion weight, execution phase
	<code>wp_exec</code>	[0.1, 0.9]	posture weight, execution phase
	<code>wg_exec</code>	[0.0, 1.0]	gaze weight, execution phase
Emotion sub-score coefficients	<code>gamma</code>	[0.0, 1.0]	valence coefficient γ
	<code>delta</code>	[0.0, 1.0]	arousal coefficient δ
Gaze cone thresholds	<code>yaw_threshold</code>	[5°, 40°]	$ \psi < \text{this} \Rightarrow \text{engaged}$
	<code>pitch_threshold</code>	[5°, 40°]	$ \phi < \text{this} \Rightarrow \text{engaged}$
Posture event detectors	<code>face_cover_ratio</code>	[0.3, 0.9]	mouth-cover fires if ratio < this
	<code>withdrawal_threshold_m</code>	[0.05, 0.30]	torso Δz for withdrawal (m)
	<code>posture_drop_gate</code>	[0.10, 0.25]	withdrawal gated by posture drop
Posture event penalties	<code>mouth_cover_penalty</code>	[0.0, 50.0]	points subtracted on mouth-cover
	<code>withdrawal_penalty</code>	[0.0, 50.0]	points subtracted on withdrawal
Temporal smoothing	<code>ema_time_constant_s</code>	[0.1, 2.0]	EMA τ , FPS-independent (s)
Missing-detection EMA decay	<code>no_face_target</code>	[30, 80]	emotion decay target, face missing
	<code>no_face_rate</code>	[0.05, 0.50]	decay rate, face missing (1/s)
	<code>no_pose_target</code>	[30, 80]	posture decay target, pose missing
	<code>no_pose_rate</code>	[0.05, 0.50]	decay rate, pose missing (1/s)

Why these groupings matter. Three of the seven groups carry most of the empirical weight. The *phase-aware fusion weights* (rows 1–6) determine which sub-score dominates in each interaction phase; the no-pose ablation discovery is literally the observation that `wp_intent` and `wp_exec` calibrate cleanly to zero under our current domain shift. The *gaze cone thresholds* widened from the default 25°/20° to roughly 32°/30° in the deployed config — the change that let `sc02`’s gaze signal stay positive while users reached for the cup. The *missing-detection EMA decay* group (rows 17–20) was hardcoded before this work; promoting these four constants to tunables was the single change that flipped the inverted `sc01/sc02` ordering, by raising `no_face_rate` from 0.15/s to roughly 0.47/s and dragging `sc01`’s sparse face-detection frames toward the low end of the comfort

range.

What is *not* in the search space. The pretrained perception models are frozen — we do not fine-tune RetinaFace, EfficientNet-B0, L2CS-Net, or MediaPipe Pose. The **approach**-phase fusion weights are also fixed at defaults ($w_e=0.5$, $w_p=0.5$, $w_g=0.3$) because the approach window contributes only a small fraction of frames per recording and adding three more dimensions for marginal gain was not justified by the information density of the dataset. The abort threshold τ^* is selected post-hoc from the test-set ROC sweep; including it inside DE would conflate calibration with operating-point selection.

Optimization Objective

The calibration searches for a parameter vector $\theta \in \mathbb{R}^{20}$ (Table 2) that minimizes the per-fold cross-validated loss

$$\mathcal{L}(\theta) = -\left[\bar{J}(\theta) + \varepsilon \cdot \frac{\mu_{\text{sc02}}(\theta) - \mu_{\text{sc01}}(\theta)}{100}\right], \quad (1)$$

with $\varepsilon = 0.1$. Equation (1) pairs a primary classification term $\bar{J}(\theta)$ with a small linear margin tie-breaker. The two terms are derived below.

Per-recording reduction. For each training recording i , the runtime pipeline is replayed under parameters θ , producing a per-frame integrated comfort series $\tilde{C}(t; \theta)$. The recording is summarized by the mean over its intent + execution window W_i (defined by the sidecar JSON labels):

$$\bar{C}_i(\theta) = \frac{1}{|W_i|} \sum_{t \in W_i} \tilde{C}(t; \theta).$$

Cross-validated Youden’s J. The training set is split into $K=5$ label-stratified folds. For each held-out fold f , the threshold τ is swept over $[30, 80]$ and Youden’s J is taken at its fold-wise maximum:

$$J_f(\theta) = \max_{\tau \in [30, 80]} \left(\frac{\text{TP}_f(\tau)}{n_f^{\text{abort}}} + \frac{\text{TN}_f(\tau)}{n_f^{\text{cont}}} - 1 \right),$$

where $\text{TP}_f(\tau)$ counts abort-labelled recordings in fold f with $\bar{C}_i \leq \tau$ and $\text{TN}_f(\tau)$ counts continue-labelled recordings with $\bar{C}_i > \tau$. The cross-validated term is the mean across folds:

$$\bar{J}(\theta) = \frac{1}{K} \sum_{f=1}^K J_f(\theta).$$

Scenario margin (tie-breaker). $\mu_{\text{sc02}}(\theta)$ and $\mu_{\text{sc01}}(\theta)$ are the means of $\bar{C}_i(\theta)$ over all training recordings labelled with each scenario, across all folds. The margin term explicitly rewards separation between the representative continue scenario (**sc02**, comfortable handoff) and the representative abort scenario (**sc01**, walk-by), independent of where the optimal threshold lands. The weight $\varepsilon = 0.1$ makes this a secondary signal; the primary optimization pressure comes from $\bar{J}$.

Optimizer. Equation (1) is minimized using SciPy’s `differential_evolution` with population multiplier `popsize=15` (yielding $15 \times 20 = 300$ candidates per generation), `maxiter=50`, convergence tolerance 10^{-3} , fixed seed 42, deferred whole-generation updating, and a final L-BFGS-B polishing pass. Differential evolution is gradient-free, which is required because $\bar{J}$ is a rank-based classification metric and non-differentiable in θ .

Variable summary.

Reading the loss in plain terms. Equation (1) says: *find parameters so that some threshold cleanly separates aborts from continues on the per-recording mean comfort (Youden’s J), and as a tie-breaker, make the comfortable-handover scenario score higher than the walk-by scenario.* Trajectory shape (whether the score trends up or down within a recording) is not in this loss — it was

Symbol	Meaning
$\theta \in \mathbb{R}^{20}$	parameter vector (Table 2)
$\tilde{C}(t; \theta)$	per-frame integrated comfort score, $\in [0, 100]$
W_i	intent + execution window of recording i (sidecar-derived)
$\bar{C}_i(\theta)$	per-recording mean of $\tilde{C}$ over W_i
τ	abort threshold, swept in $[30, 80]$ within each fold
$J_f(\theta)$	Youden’s J on fold f at the fold-optimal τ
$\bar{J}(\theta)$	mean Youden’s J across the $K=5$ CV folds
$\mu_{\text{sc02}}, \mu_{\text{sc01}}$	training-set scenario means of $\bar{C}_i$
ε	margin weight = 0.1 (tie-breaker)
K	number of stratified CV folds = 5

an evaluation gate against which the final calibrated parameters were inspected after the search. The shape-aware objective variants F, G, H, and I that replaced $\bar{J}$ with tanh-encoded slope terms all reduced classification accuracy and lost the variant race; Variant A (Eq. 1) remained the winner.